



APACHE KAFKA Tutorials/Hands-On

Kafka is a distributed , partitioned , replicated commit log service.

Fast -> Think Big Data!

-> Handle hundreds of MBs of reads & writes per second from many clients

-> Designed for real time activity streams

Distributed & highly Scalable

-> Cluster-Centric Design

-> Grow Elastically & transparent

Kafka was created to serve as a centralized online data pipelining system:

----> Kafka decouples the data pipelines.

Kafka Use-Cases

Messaging

Website Activity Tracking

Metrics

Log Aggregation

Stream Processing

Main Kafka concepts

Topics

- > Feeds of messages are organized into topics
- >Category or feed name to which messages are published
- >Each topic consists of partitioned log
 - >Partitions allow log to scale horizontally
- >Partitions are ordered,immutable sequence of messages that are continually appended
- >Each message in partitions assigned sequential id-number ("the offset")
 - >Uniquely identifies message within partition
- >Messages retained for a configurable period of time.
- >Messages consist of fixed-size header and opaque payload

Brokers

- >Kafka runs on a cluster of servers
- >Kafka Cluster comprised of one or more servers-called "brokers"
- >Each Kafka broker stores one or more partitions
 - >Can spread a single topic's partitions across multiple brokers
- >Kafka Brokers are stateless
 - >Consumers have to keep track of what offset position they would like to read

Replication

- >Each Partition gets replicated across a number of servers
- >Each partition has one "leader" server & zero or more "followers"
 - >Leader handles all read & write requests for the partition.
- >If a leader fails, a follower elected new leader.
 - >Leader handles all read & write requests for the partition
- >Followers passively replicate the leader

- >If leader fails, a follower elected new leader

- >Leaders keep track of In Sync Replicas - "ISR"

- >Message "committed" when all ISR for partition have applied it to their log
 - >Only committed messages are given to consumers
 - >Producers have option of whether to wait for commit (latency vs durability)

Producers

- >Processes that publish message to a kafka topic
- >Producer decides which message goes to which partition
 - >Round-robin for simple load balancer or custom defined.
- >Producers write to single leader
 - >Provides load balancing-different broker can service each write.
- >Kafka comes with a command-line producer client or you can write your own.

Consumers

- >Processes that subscribe to topics & process the feed of published messages.
- >Consumers read from topics
- >Kafka comes with a command line consumer client or can write your own.
- >Consumers belong to a consumer group
 - >Each message published to a topic is delivered to one consumer instance within each subscribing consumer group.
 - >Consumer instances in a consumer group can be in separate processes or separate machines.
 - >All consumer instances in the same group have load balanced across them.

->Consumers instances in different groups that subscribe to the same topic each get a copy of the messages.

Kafka requires Zookeeper

ZOOKEEPER

->Kafka requires Zookeeper to do things like

- : Cluster membership
- : Electing a controller
- : Topic configuration(which topics)

V0.8 -> Older "high-level" consumers used ZooKeeper (tracking offsets...)

V0.9+ -> Only broker uses Zookeeper

-> New consumers using instead of Zookeeper

Kafka Guarantees

->Messages sent by a producer to a particular topic partition will be appended in the order they are sent.

->Consumer instances sees messages in the

order they are stored in the log

->Committed messages will not be lost as long as at least one in sync replica alive , at all times.

Kafka Cluster Configurations

Single node, single broker

Single node , multiple brokers

Multiple nodes, multiple brokers

LAB1

Ambari-> iopbeta42.locadomain:8080/

```
#cd /usr/iop/current/kafka-broker/
```

```
#ls
```

```
#bin/kafka-topics --list --zookeeper  
localhost:2181
```

Kafka-Command-Line

Zookeeper scripts:

START

zookeeper-server-start.sh

STOP

zookeeper-server-stop.sh

**Tool used when doing migration to update
acl of znodes**

zookeeper-security-migrate.sh

Runs Zookeeper Shell

#bin/zookeeper-shell.sh localhost:2181

Kafka server start & stop

->Starts the Kafka server("broker")

->Pass in a server properties file

kafka-server-start.sh

#bin/kafka-server-start.sh

config/server.properties

kafka-server-stop.sh

Running Multiple brokers on a single machine

Generally would want to run brokers on separate machines

Each brokers gets own properties file

->

```
#cp config/server.properties  
config/server-1.properties
```

```
#cp config/server.properties  
config/server-2.properties
```

Kafka Topics

Two ways to create topics:

1.Enabling the auto.create.topics.enable
property

-> Topic created when broker receives first messages for non-existent topic

->Topic created based on num.partitions and default.replication.factor.settings

2.Using bin/kafka-topics.sh

```
bin/kafka-topics.sh --create --zookeeper  
localhost:2181 --replication-factor 1
```

```
--partitions 1 --topic mytest
```

List & describe topics

```
#bin/kafka-topics.sh --list --zookeeper  
localhost:2181
```

```
#bin/kafka-topics.sh --describe --zookeeper  
localhost:2181 --topic the-replicated-topic
```

Modifying Topics

->can increase the no of topics in a
partition

-> Pre-existing data in the topic will not
be reshuffled

```
#bin/kafka-topics.sh --zookeeper  
localhost:2181 --alter --topic my_topic  
--partition 8
```

Add configs

```
#bin/kafka-topics.sh --zookeeper  
localhost:2181 --alter--topic my_topic  
--config retention.ms=142800000
```

Remove configs

```
#bin/kafka-topics.sh --zookeeper  
localhost:2181 --alter--topic myh_topic  
--deleteConfig retention.ms
```

Deleting Topics

To enable deletion first set
delete.topic.enable=true

```
#bin/kafka-topics.sh --zookeeper  
localhost:2181 --delete --topic my_topic
```

Kafka Console Producer

Producer command line client

Start up the console producer.

```
#bin/kafka-console-producer.sh --broker-list  
localhost:9092 --topic mytopic
```

Kafka Console consumer

Consumer command line client

->Prints messages from topics to standard
output

->Each message printed on a new line

->Script runs stopped

Start up the consumer command line client

```
#bin/kafka-console-consumer.sh --zookeeper  
localhost:2181 --topic mytopic  
--from-beginning
```

Consumer groups & offset checking

```
#./kafka-consumer-groups.sh --zookeeper  
localhost:2181 --describe --group  
test-consumer-group GROUP, TOPIC , PARTITION  
, CURRENT OFFSET , LOG END OFFSET , LAG ,  
OWNER
```

LAB2

```
#cd /usr/iop/current/kafka-broker/  
ls
```

```
#bin/kafka-topics.sh --list --zookeeper  
localhost:2181
```

```
#bin/kafka-topics.sh --create --zookeeper  
localhost:2181 --replication-factor 1  
--partitions 1 --topic test_topic
```

```
#bin/kafka-topics.sh --list --zookeeper  
localhost:2181
```

```
#bin/kafka-topics.sh --create --zookeeper
localhost:2181 --replication-factor 1
--partitions 1 --config
max.message.bytes=64000 --topic
new_delete_this_topic
```

```
#bin/kafka-topics.sh --alter --zookeeper
localhost:2181 --topic delete_this_topic
-delete-config max.message.bytes
```

```
#bin/kafka-console-producer.sh --broker-list
abeta42.localdomain:6667 --topic test_topic
This is message 1
This is message 2
```

```
#cd /usr.iop/current/kafka-broker/
```

```
#bin/kafka-console-consumer.sh --zookeeper
localhost:2181 --topic test_topic
--from-beginning
```

```
#cd /usr/iop/current/kafka-broker/
```

```
#bin /kafka-topics.sh --describe --zookeeper
localhost:2181 --topic tst_topic
```

```
#pwd -> /usr/iop/current/kafka-broker
```

```
#cd config
```

```
#cat consumer.properties
```

```
#bin/kafka-console-consumer.sh --zookeeper  
localhost:2181 -topic test_to_topic  
--from-beginning --consumer.config  
config/consumer.properties.
```

```
#bin/kafka-consumer-groups.sh --describe  
--zookeeper localhost:2181 --group  
test-consumer-group
```

KAFKA Producer Client Overview

- >Kafka includes a new producer client
- >Java API easy to work with.
- >Variety of other source producers also available
 - >Python
 - >C++
 - >C#
- >New Java producer replaces older Scala producer client (V0.8.2)
- >Use the API to code your own producer logic
- >Thread Safe producer sends data directly to

the broker that is the leader for the partition

- >Maps messages to topic partition

- >Sends produce requests to leader of that partition

- >Client controls which partition it publishes messages to

 - >Hash partitioning if user specifies key to partition by

 - >Random load balancing if no key provided (round-robin)

- >Batching is big driver of efficiency

 - >Producer will attempt to accumulate data in memory and to send out larger batched in a single request

 - >Batching is configurable.

 - >Trades off small amount of latency for better throughput

Producer configuration Settings

Variety of producer configuration settings

- >**bootstrap.servers**

 - >List of host/port pairs used for establishing initial connection to kafka cluster

->client.id

->id string passed to the server when making requests

->Makes tracking source of requests easier for logging vs using ip/port

->acks

->Number of acknowledgments producer requires leader to have received before considering a request complete - Controls durability of sent records

- acks=0 : producer will not wait for any acknowledgement from the server at all

- acks=1 : leader writes record to local log but responds without awaiting full acknowledgement from all followers

- acks=all : leader waits for full set of in-sync replicas (ISR) to acknowledge record

->retries

->Value >0 causes to resend any record whose send fails with a potentially transient error

->Allowing retries will potentially change the ordering of records

->batch.size

->Producer attempts to collect messages into batches.

->Increases throughput

->buffer.memory

->Limit total memory used to store each batch

->linger.ms

->Have producer delay sending to give more time for batches to be filled

->compression.type

->Compression type for all data generated by the producer

->Valid values: none (default), gzip , snappy, or lz4.

->Compression is of full batches of data

->key.serializer and value.serializer

->How to turn key & value objects user provides with ProducerRecord into bytes.

Example of Configuration code:

```
Properties props = new Properties();
```

```
props.put("bootstrap.servers","localhost:9092");
props.put("acks","all");

props.put("retries",0);

props.put("batch.size",16384);

props.put("linger.ms",1);

props.put("buffer.memory",33554432);

props.put("key.serializer","org.apache.kafka.
common.serialization.StringSerializer");

props.put("value.serializer","org.apache.kafk
a.common.serialization.StringSerializer");
```

KafkaProducer Java Class Overview:

- >Class KafkaProducer<K,V>
- >Kafka Client that publishes records to the Kafka Cluster
- >Producer is thread safe.
- >Producer consist of :
 - >Pool of buffer space that holds records that haven't yet been transmitted to server.

->Background I/O thread responsible for turning records into requests and transmitting them to cluster

Methods:

```
close();  
close(long  
timeout,java.util.concurrent.TimeUnit  
timeUnit);  
flush();  
metrics();  
partitionsFor(java.lang.Stringtopic)  
send(ProducerRecord<K,V>record)  
send(ProducerRecord<K,V>record , Callback  
callback)
```

KafkaProducer constructor:

->KafkaProducer(java.util.Map<java.lang.String,java.lang.Object> configs)

-> A producer is instantiated by providing a set of key-value pairs as configuration.

->KafkaProducer(java.util.Map<java.lang.String,java.lang.Object> configs, Serializer<K> keySerializer , Serializer<V>

valueSerializer)

-> A producer is instantiated by providing a set of key-value pairs as configuration. a key & a value serializer

->KafkaProducer(java.util.Properties properties)

-> A producer is instantiated by providing a set of key-value pairs as configuration.

->KafkaProducer(jav.util.Properties properties , Serialzer<K> keySerializer , Serializer<V> valueSerializer)

-> A producer is instantiated by providing a set of key-value pairs as configuration , a key & a value Serializer.

Example KafkaProducer constructor code:

```
Properties props = new Properties();
props.put("bootstrap.servers","localhost:9092");
props.put("acks","all");
props.put("key.serializer","org.kafka.common.serialization.StringSeriliazer");
props.put("value.serailizer","org.apache.kafk
```

```
a.common.serialization.StringSerializer");
```

```
KafkaProducer<String, String> producer  
    = new KafkaProducer<String, String(props)>;
```

KafkaProducer send method:

```
-> public  
java.util.concurrent.Future<RecordMetadata> send  
(ProducerRecord<K,V> record , Callback callback)
```

-> All writes are asynchronous by default.

->Method returns immediately once record has been stored in a buffer of records waiting to be sent.

->Allows sending many records in parallel without blocking to wait for a response after each one.

-> Send() API returns a Future for RecordMetadata which can be polled to get result of the send

->Invokes the provided callback when the send has been acknowledged.

->Result of send is RecordMetadata specifying partition & offset.

->Invoking get() on future, blocks until associated request completes

->Then return metadata for record or throws any exception that occurred while sending the record.

->For synchronous writes , use get() immediately to simulate simple blocking call

ProducerRecord and RecordMetadata

Class ProducerRecord

ProducerRecord(java.util.String topic ,
java.lang.Integer partition, K key , V value)

->A key/value pair(message) to be sent to kafka.

->Consists of topic name to which record is being sent , optional partition number , and optional key ,and the message value.

->If valid partition number is specified - that partition is used when sending a record.

->If no partition specified but key is present , partition chosen using hash of key

->If neither key nor partition present a partition assigned in round-robin fashion.

Class RecordMetadata

->The metadata for a record that has been acknowledged by the server

->Send() API returns a Future for this RecordMetadata

->Can use this to get offset , topic , partition

Asynchronous KafkaProducer send example:

```
ProducerRecord<byte[],byte[]> myRecord = new  
ProducerRecord<byte[],byte[]>("mytopic",key,value);
```

```
producer.send(myRecord , new Callback()  
{  
    public void onCompletion(RecordMetadata metadata  
,Exception e)  
    {  
        if(e!=null)  
            e.printStackTrace();  
        System.out.println("Offset of sent record:"  
+metadata.offset());  
    }  
});
```

Kafka Consumer Client Overview

->Kafka includes a new consumer client that consumes records from the Kafka cluster.

->Java API easy to work with:

Class `kafkaConsumer<k,V>`

->Variety of other source producers also available

->Python

->C++

->C#

->Use the API to code your own consumer logic.

->Transparently handles failure of servers in kafka cluster

- >Interacts with the server to allow groups of consumers to load balance consumption using consumer groups.
- >Maintains TCP connections to the necessary brokers to fetch data.
- >You can let consumer automatically handle much of the fine-grained details (like prior "high-level" consumer) OR you can manage lower-level features (like prior "simple" consumer)

General consumer application flow

- >Create consumer configuration & use to create new consumer instance.
- >Subscribe to topics/partitions
- >Poll the topic(s) from within a loop to fetch messages
- >Do something useful with the messages consumed
- >Commit offsets(automatically or manually)
- >Close the consumer.

Group coordination

->Consumers use group coordination protocol built into Kafka (prior version used zookeeper)

->For each group, one of the brokers is selected as the group coordinator.

->Coordinator responsible for managing the state of the group.

->Act of reassigning partitions is called rebalancing the group.

->Rebalancing is process where group of consumer instances within a consumer group, Coordinate to own a mutually exclusive set of partitions of topics that the group has subscribed to.

->Coordinator monitors the heartbeat of group members to ensure they are alive.

Manual vs Automatic Offset committing

->Consumer can automatically commit offsets, or you can manually commit

->Config **enable.auto.commit**

->enable.auto.commit=true:

Offsets are committed automatically with frequency controlled by auto.commit.interval.msconfig

->enable.auto.commit=false :

Manually offset commits

->Committing offsets manually allows you to perform processing logic and only commit once processing is completed.

->For Example: batch a quantity of records & then submit to a database in a single transaction. Commit after the transaction is complete.

->Provides "at-least-once delivery" guarantees

->If auto-committing in this scenario, could commit after messages consumed but before processing logic completes resulting in lost messages and at-most-once delivery guarantee.

->Manually commit offsets with commitSync() and commitAsync() methods

Consumer configuration settings

Variety of consumers configuration settings

->bootstrap.servers

->List of host/port pairs used for establishing initial connection to kafka cluster

->client.id

->id string passed to the server when making requests

->group.id

->Unique string that identifies consumer groups that consumer belongs to

->key.deserializer and value.deserializer

->How to deserialize message keys and values

->fetch.min.bytes

->Minimum amount of data the server should return for a fetch request

->1 byte by default - means fetch request answered as soon as single byte of data is available

->enable.auto.commit

->Set to true for automaticcommits

->auto.commit.interval.ms

->The frequency in milliseconds that the consumer offsets are auto-committed to kafka if `enable.auto.commit` is set to true.

->`session.timeout.ms`

->Timeout used to detect failures of a consumer

->Kafka group coordinator expects to receive a heartbeat from consumers within this time period.

->`heartbeat.interval.ms`

->Expected time between heartbeats to the consumer coordinator

->Heartbeats used to ensure consumers' sessions stay active & to facilitate rebalancing when new consumers join or leave the group.

Configuration CODE example

```
Properties props = new Properties();
props.put("bootstrap.servers","localhost:9092");
props.put("group.id","testgroup");
props.put("enable.auto.commit","true");
props.put("auto.commit.interval.ms","1000");
props.put("session.timeout.ms","30000");
```

```
props.put("key.deserializer","org.apache.kafka.common.serialization.StringDeserializer");
```

```
props.put("value.deserializer","org.apache.kafka.common.serialization.StringDeserializer");
```

KafkaConsumer Java Class Overview:

KafkaConsumer constructor

->KafkaConsumer(java.util.Map<java.lang.String, java.lang.Object> configs)

->Consumer instantiated by providing set of key-value pairs as configuration

->KafkaConsumer(java.util.Map<java.lang.String, java.lang.Object> configs, Deserializer<K> keyDeserializer , Deserializer<V> valueDeserializer)

->Consumer instantiated by providing set of key-value pairs as configuration , a ConsumerRebalanceListener implementation , a key & value Deserializer

->KafkaConsumer(java.util.Properties properties)

->Consumer instantiated by providing object as configuration

->KafkaConsumer(java.util.Properties properties, Deserializer<K> keyDeserializer , Deserializer)

->Consumer instantiated by providing Properties as objects as configuration and a ConsumerRebalanceListener implementation , a key & value Deserializer.

Subscribing to Topics

->Subscribe to the given list of topics to get dynamically assigned partitions

->Can subscribe with a list of topics or a regex pattern.

->Kafka assigns fair share of partitions from these topics to consumers

->Topics subscriptions are not incremental

->List replaces the current assignment (if there is one)

```
KafkaConsumer<String,String> consumer = new  
KafkaConsumer<>(props);
```

```
consumer.subscribe(Arrays.asList("topic1","to  
pic2"));
```

Manual partition assignment

->can alternatively subscribe to specific partitions.

->Consumer will just get the partitions it subscribes to

->If consumer instance fails no attempt is made to rebalance partitions

```
TopicPartition partition0 = new  
TopicPartition("mytopic",0);
```

```
TopicPartition partition1 = new  
TopicPartition("mytopic",1);
```

```
consumer.assign(Arrays.asList(partition0,partition1));
```

Polling

```
public ConsumerRecords<K,V> poll(long  
timeout)
```

-> Use poll to fetch messages that consumers is subscribed to

->timeout

- >Time in millis ,spent waiting in poll if data is not available
- >If 0 , returns immediately with any records that are available now
- >Returns map of records since last fetch for subscribed list of topics and partitions
- >Meant to be run in an event loop

```
try{
    while(running)
    {
        ConsumerRecords<String,String> records
= consumer.poll(100) ;
        //Process the records...
    }
}
```

Close the Consumer

```
public void close()
```

- >Use to close the consumer

- >Waits indefinitely for any needed cleanup

->If auto-commit enabled calling close() will commit current offsets'

->Always call close()

->Cleans up any sockets in use

->Ensures consumer can alert coordinator about its departure from the group

```
try{
    while (running) {
        ConsumerRecords <String,String> records =
consumer.poll(100);
    }
}finally{
    consumer.close();
}
```

Storing offsets outside of kafka & controlling consumer's position

->Can choose to store offsets outside of kafka

->Allows for fully atomic consumption and "exactly-one" semantics
->Stronger than "at-least-once" semantics that come with kafka offset commit functionality

->For Ex. could store offsets along with processed data at same time in a database transaction

->Either transaction succeeds and offset updated or transaction fails and neither offset nor processed data is stored

->seek(TopicPartition partition ,long offset)

->seekToBeginning(TopicPartition .. partitions)

->seekToEnd(TopicPartition ...partitions)

For Example

,could use seek to re-consume messages if transaction failed.

```
javac -d MyKafkaConsumer.java ->  
java MyKafkaConsumer
```

```
#pwd
```

```
#bin/kafka-console-producer.sh --broker-list  
abeta42.localdomain:6667 --topic test_topic  
This is a test 1
```

```
#try with manual offsets
```

Kafka Connect & Spark streaming

Kafka Connect

**->Adopting Kafka for data integration
required significant development skills**

->Developing a kafka connector required
building on the client API's

->New(V0.9)Kafka Connect is framework for
large scale , real-time stream data
integration using kafka.

->Simplifies adoption of connectors for
stream data integration

->Makes building and managing stream data
integration.

->Encourages development of rich
ecosystems of open source connectors

->Deploy Kafka connectors that work well with each other and can be monitored, deployed , and administered in consistent manner

->Focusses on the ETL

->Kafka must always be on one side of the equation - source or sink.

->Abstracts away common problems every connector to Kafka faces:

->Schema management , fault tolerance , partitioning , offset management & delivery semantics & monitoring

Kafka connect features

->A common framework for kafka connectors

->Distributed & standalone modes

->REST interface

->Automatic offset management

->Streaming/batch integration

->Leveraging Kafka's existing capabilities ,
Kafka Connect great for bridging streaming
and batch data systems

Connectors

->Can develop your own connectors or use an
already available one
->Some nice pre-built connectors out there
 ->HDFS Connector
 ->Exports Data from Kafka topics to HDFS
files in variety of formats
 ->Periodically polls data from Kafka and
writes them to HDFS
 ->Data from each Kafka topic is
partitioned and divided into chunks
 ->Each chunk of data is represented as an
HDFS Filewith topic, Kafka partition, start
and end offsets of these data chunks in the
filename.

JDBC Connector

->Import data from any relational database
with JDBC driver into Kafka Topics.

->Data is loaded by periodically executing a SQL Query and creating an output record for each row in the result set.

->The database is monitored for new or deleted tables and adapts automatically.

->FileStream Connector

->Variety of others being developed

->MongoDB, Cassandra

Connectors , Tasks and Workers

->Connector instance

->Logical job responsible for managing copying of data between Kafka & other systems.

->Connector either source or sink connector

->Each Connector can instantiate a set of tasks that actually copy the data.

->Ability to break a single job to multiple tasks allows for parallelism.

->All Kafka Connect sources and sinks map to partitioned streams of records

->A Kafka Connect cluster consists of set of Worker Process

->ex-if streams represents a database , each stream partition would represent table in the DB

->Every record in partitioned stream consists a key, a value , and associated offset

->Offsets are tracked by kafka Connect and mark position of every record in the stream partition.

->Similar in concept to kafka offsets , however can be different formats

->For Example offsets for records coming from database might be timestamp column.

->A Kafka Connect cluster consists of set of Worker processes

->Containers that execute Connectors and Tasks.

->Auto coordinate with each other to distribute work and provide scalability and fault tolerance

->Connect has standalone and distributed workers.

Worker Standalone vs Distributed modes

Standalone

- >Simplest mode
- >Single process responsible for executing all connectors and tasks
- >Requires minimal configuration -passwd in via command line

```
bin/connect-standalone.sh worker.properties
```

```
Connector1.properties[connector2.properties.]
```

Distributed

- >Provides scalability and automatic fault tolerance
- >Have many worker process share a group.id
 - >Workers will automatically coordinate to schedule execution of connectors and tasks across all available workers
- >Interaction with distributed-mode cluster using REST API

->Example: GET /connectors - returns a list of active connectors

Kafka Connect StandAlone Example

#bin/connect-standalone.sh
config/connect-standalone.properties
config/connect-file-source.properties

Kafka & Spark

->Many use cases entail Kafka feeding streaming data into spark streaming
->Newer Spark Streaming/Kafka Direct API (introduced Spark 3.1)

Kafka => Spark Streaming
=>HDFS/Databases/DASHboards